

Refactoring

Code

Setup & Imports

```
[ ] # -*- coding: utf-8 -*-  
    """  
    Created on Thu Nov 21 17:51:09 2024  
  
    @author: Sheila  
    """  
  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
  
from sklearn.impute import SimpleImputer  
from sklearn import preprocessing  
from sklearn.model_selection import train_test_split  
  
from sklearn.naive_bayes import GaussianNB  
from sklearn.neighbors import KNeighborsClassifier  
from sklearn.tree import DecisionTreeClassifier  
from sklearn.ensemble import RandomForestClassifier  
from sklearn.svm import SVC  
from xgboost import XGBClassifier  
from sklearn.linear_model import LogisticRegression  
from sklearn.ensemble import AdaBoostClassifier  
  
from sklearn.metrics import confusion_matrix  
from sklearn.metrics import accuracy_score, classification_report, auc  
from sklearn.metrics import roc_curve
```

Load and Explore Dataset

```
[2] from google.colab import files
     uploaded = files.upload()
```



Choose Files heart_statlo...ry_final.csv

- **heart_statlog_cleveland_hungary_final.csv**(text/csv) - 38805 bytes, last modified: 11/27/2024 - 100% done
- Saving heart_statlog_cleveland_hungary_final.csv to heart_statlog_cleveland_hungary_final.csv

```
[3] # Load dataset
     df = pd.read_csv("heart_statlog_cleveland_hungary_final.csv")

     df.head()
     df.info()
     df.describe()

     #to count how many have and do not have heart disease
     df['target'].value_counts()

     #To get the percentage distribution of values of a column we use df['colname'].value_counts(normalize=True)*100
     df['target'].value_counts(normalize=True)*100

     # Plot target distribution
     #sns.catplot(x="target", kind="count", hue="target", palette="ch:.25", legend=False, data=df).set(title="Distribution of normal (0) vs with heart disease (1)")
     sns.catplot(x="target", kind="count", hue="target", palette="ch:.25", legend=False, data=df).set(title="Distribution of normal (0) vs with heart disease (1)")
```

```
[ ] print(" = Data Visualization = ")
# box and whisker plots
print(" == Univariate Plots: box and whisker plots. why? to determine outliers? = ")
df.plot(kind='box', subplots=True, layout=(4,3), sharex=False, sharey=False, figsize=(10,8))
plt.show()

# Correlation map
print(" == Correlation Map == ")
print("Here we see that resting bp s and resting ecg have correlation < 0.2 with the outcome hence are candidates for removal")
plt.figure(figsize=(10,10))
sns.heatmap(df.corr(), annot=True, square=True, cmap='magma')
plt.xticks(rotation=60)
plt.show()

df.isnull().any()
```

Data visualization

Changes made

- Code was put inside a function
 - `visualize_data(df, target_column='target', corr_threshold=0.2)`
- No more hardcoded values
- Correlation threshold and target can be provided during a function call
- Automatically prints features with < 0.2 (weak) correlation with target

Data Visualization

Refactored Code

```
[ ] visualize_data(df, target_column='target', corr_threshold=0.2)
```

```
[ ] def visualize_data(df, target_column=None, corr_threshold=0.2, box_layout=(4, 3), box_figsize=(10, 8), heatmap_figsize=(10, 10), heatmap_cmap='magma'):  
    """  
    Visualizes data using box and whisker plots and a heatmap, with dynamic correlation analysis.  
  
    Parameters:  
    - df: pandas DataFrame  
    - target_column: str or None, column name for the target variable (if correlation analysis is required)  
    - corr_threshold: float, threshold below which features are considered weakly correlated with the target  
    - box_layout: tuple, layout for the boxplot grid (rows, columns)  
    - box_figsize: tuple, figure size for the boxplot  
    - heatmap_figsize: tuple, figure size for the heatmap  
    - heatmap_cmap: str, colormap for the heatmap  
    """  
    print(" = Data Visualization = ")  
  
    # Box and Whisker Plots  
    print(" == Univariate Plots: Box and Whisker Plots ==")  
    print("Purpose: To identify outliers and understand data distribution.")  
    df.plot(kind='box', subplots=True, layout=box_layout, sharex=False, sharey=False, figsize=box_figsize)  
    plt.tight_layout()  
    plt.show()  
  
    # Correlation Heatmap  
    print(" == Correlation Map ==")  
    print("Purpose: To identify potential feature relationships.")  
  
    plt.figure(figsize=heatmap_figsize)  
    sns.heatmap(df.corr(), annot=True, square=True, cmap=heatmap_cmap)  
    plt.xticks(rotation=60)  
    plt.tight_layout()  
    plt.show()
```

Data Visualization

Refactored Code

```
# Dynamic Weak Correlation Analysis
if target_column and target_column in df.columns:
    correlations = df.corr()[target_column].round(2)
    weak_corr_features = correlations[correlations.abs() < corr_threshold].index.tolist()
    if weak_corr_features:
        print(f"Features weakly correlated with '{target_column}' (|correlation| < {corr_threshold}): {weak_corr_features}")
    else:
        print(f"All features have a strong correlation (|correlation| >= {corr_threshold}) with '{target_column}'")
else:
    print("No target column provided or column not found in the dataset. Skipping correlation threshold analysis.")

df.isnull().any()
```

```
# Selecting duplicate rows except first
# occurrence based on all columns
# saves the duplicates in the dataframe called dup
dup = df[df.duplicated()]
print("These are the duplicate rows :")
dup
dup.shape

# save a copy of the records with no duplicates in a new dataframe named df_nodup
df_nodup = df.drop_duplicates()
df_nodup.to_csv('df_nodup.csv')
df_nodup.shape

df.duplicated()

print('Shape of original dataframe from csv: ', df.shape)
print('Shape of dataframe containing the duplicates: ', dup.shape)
print('Shape of dataframe with no more duplicates: ', df_nodup.shape)
```

Rerun EDA on the dataframe with no duplicates

Original Code

```
#to count how many have and do not have heart disease
df_nodup['target'].value_counts()

#To get the percentage distribution of values of a column we use df['colname'].value_counts(normalize=True)*100
df_nodup['target'].value_counts(normalize=True)*100

#sns.catplot(x="target", kind="count", palette="ch:.25", data=df_nodup).set(title="Distribution of normal (0) vs with heart disease (1)")
sns.catplot(x="target", kind="count", hue="target", palette="ch:.25", legend=False, data=df_nodup).set(title="Distribution of normal (0) vs with heart disease (1)")
```


Rerun EDA on the dataframe with no duplicates

Refactored Code

Changes made

- Exact number of heart disease counts is displayed
- Percentage is displayed
- Labels were added to the graph

```
# Count the number of individuals with and without heart disease
heart_disease_counts = df_nodup['target'].value_counts()
print("Heart Disease Counts:")
print(heart_disease_counts)

# Calculate the percentage distribution of heart disease status
heart_disease_percentage = df_nodup['target'].value_counts(normalize=True) * 100
print("\nHeart Disease Percentage Distribution:")
print(heart_disease_percentage)

# Visualize the distribution of individuals with and without heart disease
sns.catplot(
    x="target",
    kind="count",
    hue="target",
    palette="ch:.25",
    legend=False,
    data=df_nodup
).set(title="Distribution of Normal (0) vs With Heart Disease (1)")

plt.legend(title="Heart Disease Status", labels=["Normal (0)", "With Heart Disease (1)"])
plt.show()
```

Rerun EDA on the dataframe with no duplicates

Output

Heart Disease Counts:

target

1 508

0 410

Name: count, dtype: int64


Heart Disease Percentage Distribution:

target

1 55.337691

0 44.662309

Name: proportion, dtype: float64



Handling Missing Values

Original Code

```
#convert missing values represented by 0's into NAs
print(df_nodup.isnull().sum())
print("These columns have zero values which are unbelievable and thus should be replaced by NaN's")
temp_cols = ['cholesterol', 'ST slope']
df_nodup.loc[:, temp_cols] = df_nodup[temp_cols].replace(0, np.nan)
print(df_nodup.isnull().sum())
print(df_nodup[['cholesterol', 'ST slope']].describe())
```

Handling missing values

Refactored Code

Changes made

- Extract method
 - Turning the method into subfunctions

```
def check_missing_values(df):  
    """Prints the count of missing values for each column."""  
    print("Missing Values in Each Column:")  
    print(df.isnull().sum())  
    print("\n")  
  
def replace_zeros_with_nan(df, columns):  
    """Replaces zero values with NaNs in specified columns."""  
    print(f"Replacing zero values with NaNs in the following columns: {columns}")  
    df[columns] = df[columns].replace(0, np.nan)  
    print("Replacement Complete.\n")  
    return df  
  
def display_column_statistics(df, columns):  
    """Displays summary statistics for specified columns."""  
    print("Summary Statistics for Specified Columns:")  
    print(df[columns].describe())  
    print("\n")
```

Handling missing values

Refactored Code

Changes made

- Extract method
 - Turning the method into subfunctions

Missing Values in Each Column:

```
age          0
sex          0
chest pain type  0
resting bp s  0
cholesterol  172
fasting blood sugar  0
resting ecg   0
max heart rate  0
exercise angina  0
oldpeak       0
ST slope      1
target        0
dtype: int64
```

Replacing zero values with NaNs in the following columns: ['cholesterol', 'ST slope']
Replacement Complete.

Missing Values in Each Column:

```
age          0
sex          0
chest pain type  0
resting bp s  0
cholesterol  172
fasting blood sugar  0
resting ecg   0
max heart rate  0
exercise angina  0
oldpeak       0
ST slope      1
target        0
dtype: int64
```

Summary Statistics for Specified Columns:

	cholesterol	ST slope
count	746.000000	917.000000
mean	244.635389	1.63795
std	59.153524	0.60727
min	85.000000	1.00000
25%	207.250000	1.00000
50%	237.000000	2.00000
75%	275.000000	2.00000
max	603.000000	3.00000

Imputing the NaN's in Insulin, Skin Thickness, BMI, Blood Pressure, and Glucose using the median

Original Code

```
[ ] # Define X and y
XMedian = df_nodup[df_nodup.columns[:-1]]
yMedian = df_nodup[df_nodup.columns[-1]]

# Print total missing before imputation
print('----- See Original Count, Mean, and STD -----\\n')
print(XMedian[['cholesterol', 'ST slope']].describe())

# Define imputer
imputer = SimpleImputer(strategy='median')

# Fit on the dataset
imputer.fit(XMedian)

# Transform the dataset
XtransMedian = imputer.transform(XMedian)

# Print total missing after imputation
print('\\n----- See Changes in Count, Median, and STD -----\\n')
print(pd.DataFrame(XtransMedian, columns=XMedian.columns)[['age', 'sex', 'chest pain type', 'resting bp s', 'cholesterol', 'fasting blood sugar', 'resting ecg', 'max heart rate', 'exercise angi

XMedian
XtransMedian = pd.DataFrame(XtransMedian, columns = ['age', 'sex', 'chest pain type', 'resting bp s', 'cholesterol', 'fasting blood sugar', 'resting ecg', 'max heart rate', 'exercise angina', '

XtransMedian

print(XtransMedian[['cholesterol', 'ST slope']].describe())
```


Imputing the NaN's in Insulin, Skin Thickness, BMI, Blood Pressure, and Glucose using the median

Refactored Code

Changes made

- Extract method
 - Turning the method into subfunctions

```
def describe_columns(dataframe, columns, message):  
    """  
    Prints descriptive statistics for the specified columns of a DataFrame.  
    """  
    print(f"----- {message} -----\n")  
    print(dataframe[columns].describe())  
    print("\n")  
  
def apply_median_imputation(dataframe):  
    """  
    Applies median imputation to the DataFrame and returns the transformed DataFrame.  
    """  
    imputer = SimpleImputer(strategy='median')  
    imputed_data = imputer.fit_transform(dataframe)  
    return pd.DataFrame(imputed_data, columns=dataframe.columns)
```

Imputing the NaN's in Insulin, Skin Thickness, BMI, Blood Pressure, and Glucose using the median

Refactored Code

Changes made

- Printed the step by step process instead of directly updating the full transformed dataframe statistics:
 - Original Count, Median, and STD
 - Updated Count, Median, and STD
 - Full Transformed DataFrame Statistics
 - Transformed DataFrame (Sample Rows)

```
def main_imputation_process(df):  
    """  
    Orchestrates the process of describing, imputing, and displaying results.  
    """  
  
    # Define feature columns (X) and target column (y)  
    X = df.iloc[:, :-1]  
    y = df.iloc[:, -1]  
  
    # Columns to inspect  
    inspect_columns = ['cholesterol', 'ST slope']  
  
    # Display original statistics  
    describe_columns(X, inspect_columns, "Original Count, Median, and STD")  
  
    # Apply median imputation  
    X_imputed = apply_median_imputation(X)  
  
    # Display updated statistics  
    describe_columns(X_imputed, inspect_columns, "Updated Count, Median, and STD After Imputation")  
  
    # Display the entire DataFrame summary to compare structure  
    describe_columns(  
        X_imputed,  
        X_imputed.columns,  
        "Transformed DataFrame Statistics"  
    )  
  
    return X_imputed  
  
# Apply the process to the dataset  
X_imputed = main_imputation_process(df_nodup)  
  
# Display the transformed DataFrame  
print("----- Transformed DataFrame (Sample Rows) -----\\n")  
print(X_imputed.head())
```


Imputing the NaN's in Insulin, Skin Thickness, BMI, Blood Pressure, and Glucose using the median

Refactored Code

----- Original Count, Median, and STD -----

	cholesterol	ST slope
count	746.000000	917.000000
mean	244.635389	1.63795
std	59.153524	0.60727
min	85.000000	1.00000
25%	207.250000	1.00000
50%	237.000000	2.00000
75%	275.000000	2.00000
max	603.000000	3.00000

----- Updated Count, Median, and STD After Imputation -----

	cholesterol	ST slope
count	918.000000	918.000000
mean	243.204793	1.638344
std	53.401297	0.607056
min	85.000000	1.00000
25%	214.000000	1.00000
50%	237.000000	2.00000
75%	267.000000	2.00000
max	603.000000	3.00000

----- Transformed DataFrame Statistics -----

	age	sex	chest pain type	resting bp s	cholesterol \
count	918.000000	918.000000	918.000000	918.000000	918.000000
mean	53.510893	0.789760	3.251634	132.396514	243.204793
std	9.432617	0.407701	0.931031	18.514154	53.401297
min	28.000000	0.000000	1.000000	0.000000	85.000000
25%	47.000000	1.000000	3.000000	120.000000	214.000000
50%	54.000000	1.000000	4.000000	130.000000	237.000000
75%	60.000000	1.000000	4.000000	140.000000	267.000000
max	77.000000	1.000000	4.000000	200.000000	603.000000

Imputing the NaN's in Insulin, Skin Thickness, BMI, Blood Pressure, and Glucose using the median

Refactored Code

```
fasting blood sugar  resting ecg  max heart rate  exercise angina  \
count      918.000000    918.000000    918.000000    918.000000
mean        0.233115    0.603486    136.809368    0.404139
std         0.423046    0.805968    25.460334    0.490992
min         0.000000    0.000000    60.000000    0.000000
25%         0.000000    0.000000    120.000000    0.000000
50%         0.000000    0.000000    138.000000    0.000000
75%         0.000000    1.000000    156.000000    1.000000
max         1.000000    2.000000    202.000000    1.000000

oldpeak  ST slope
count    918.000000  918.000000
mean     0.887364    1.638344
std      1.066570    0.607056
min     -2.600000    1.000000
25%      0.000000    1.000000
50%      0.600000    2.000000
75%      1.500000    2.000000
max       6.200000    3.000000
```

Preprocessing categorical variables using get_dummies() function of pandas

Original Code

```
#categorical data
categorical_cols = ['chest pain type', 'resting ecg', 'ST slope']

#import pandas as pd
df_dummyfied = pd.get_dummies(df_nodup, columns = categorical_cols, drop_first=True)
df_dum = df_dummyfied.drop(['age', 'resting bp s', 'cholesterol', 'max heart rate', 'oldpeak'], axis = 1)
df_dum.to_csv('df_dum.csv')

colnamesko = list(df_dum.columns)
column_names = ['sex', 'fasting blood sugar', 'exercise angina', 'chest pain type_2', 'chest pain type_3', 'chest pain type_4', 'resting ecg_1', 'resting ecg_2', 'ST slope_2.0', 'ST slope_3.0',
df_dum = df_dum.reindex(columns=column_names)
```

Preprocessing categorical variables using `get_dummies()` function of pandas

Refactored Code

Changes made

- Refactored variables
 - Preprocessing categorical variables using `get_dummies()` function of pandas

```
numerical_cols = ['age', 'resting bp s', 'cholesterol', 'max heart rate', 'oldpeak']  
categorical_cols = ['chest pain type', 'resting ecg', 'ST slope']
```

Preprocessing categorical variables using `get_dummies()` function of pandas

Refactored Code

Changes made

- Replace Magic Numbers with Constants
 - Replace ['age', 'resting bp s', 'cholesterol', 'max heart rate', 'oldpeak'] to variable `numerical_cols`
 - Replace ['chest pain type', 'resting ecg', 'ST slope'] to variable `categorical_cols`
- Consolidate duplicate code
 - Used a function to implement processing of categorical data into one

```
def process_categorical_data(df, categorical_cols, drop_columns, column_order):
    df_dummyfied = pd.get_dummies(df, columns=categorical_cols, drop_first=True)
    df_dum = df_dummyfied.drop(drop_columns, axis=1)
    df_dum.to_csv('df_dum.csv', index=False)
    df_dum = df_dum.reindex(columns=column_order)

    return df_dum

column_names = ['sex', 'fasting blood sugar', 'exercise angina', 'chest pain type_2', 'chest pain type_3',
                'chest pain type_4', 'resting ecg_1', 'resting ecg_2', 'ST slope_2.0', 'ST slope_3.0', 'target']

df_dummyfied = pd.get_dummies(df, columns=categorical_cols, drop_first=True)
df_dum = process_categorical_data(df_nodup, categorical_cols, numerical_cols, column_names)
```

Preprocessing Numeric variables using StandardScaler

Changes made

- Replace Magic Numbers with Constants
 - Replace ['age', 'resting bp s', 'cholesterol', 'max heart rate', 'oldpeak'] to variable numerical_cols
- Consolidate Duplicate Code
 - Used a function to implement Preprocessing, etc into one

```
def standardize_and_concat(df, numerical_cols):  
    std_scale = preprocessing.StandardScaler().fit(df[numerical_cols])  
    data_std = std_scale.transform(df[numerical_cols])  
  
    data_std = pd.DataFrame(data_std, columns=numerical_cols)  
    df_dum.reset_index(drop=True, inplace=True)  
    data_std.reset_index(drop=True, inplace=True)  
  
    new_df = pd.concat([data_std, df_dum], axis=1)  
  
    return new_df  
  
data_std = standardize_and_concat(df_dummyfied, numerical_cols)  
XtransMedian = standardize_and_concat(XtransMedian, numerical_cols)
```

Performance evaluation

Original Code

```
yMedian=XtransMedian['target']
xMedianMedian=XtransMedian.drop(['target'], axis=1)
XtransMedian.drop(['target'], axis=1)
x_train,x_test,y_train,y_test=train_test_split(xMedianMedian, yMedian,test_size=0.30,random_state=14)

dt = DecisionTreeClassifier(random_state=14)
dt.fit(x_train,y_train)
y_pred = dt.predict(x_test)
print("Decision Tree Accuracy with Median Imputation on 10 predictors", accuracy_score(y_test, y_pred))
print("Confusion Matrix: \n", confusion_matrix(y_test, y_pred))
print("Classification Report Decision Tree with Median Imputation on 10 predictors\n", classification_report(y_test, y_pred))
probsdt = dt.predict_proba(x_test)
fprdt, tprdt, thresholds = roc_curve(y_test, probsdt[:, 1], pos_label=1)
roc_aucdt = auc(fprdt, tprdt)

lr=LogisticRegression(C=50)
lr.fit(x_train, y_train)
y_pred=lr.predict(x_test)
print("Logistic Regression Accuracy with Median Imputation and with Pregnancy, Insulin, and Skin Thickness Removed", accuracy_score(y_test, y_pred))
print("Confusion Matrix: \n", confusion_matrix(y_test, y_pred))
print("Classification Report Logistic Regression with Median Imputation and with Pregnancy, Insulin, and Skin Thickness Removed\n", classification_report(y_test, y_pred))
probslr = lr.predict_proba(x_test)
fprlr, tprlr, thresholds = roc_curve(y_test, probslr[:, 1], pos_label=1)
roc_auclr = auc(fprlr, tprlr)

rf=RandomForestClassifier()
rf.fit(x_train, y_train)
y_pred=rf.predict(x_test)
print("Random Forest Accuracy with Median Imputation and with Pregnancy, Insulin, and Skin Thickness Removed", accuracy_score(y_test, y_pred))
print("Confusion Matrix: \n", confusion_matrix(y_test, y_pred))
print("Classification Report Random Forest with Median Imputation and with Pregnancy, Insulin, and Skin Thickness Removed\n", classification_report(y_test, y_pred))
probsrf = rf.predict_proba(x_test)
fprrf, tprrf, thresholds = roc_curve(y_test, probsrf[:, 1], pos_label=1)
roc_aucrf = auc(fprrf, tprrf)
```


Performance evaluation

Original Code

```
nb = GaussianNB()
nb.fit(x_train, y_train)
y_pred = nb.predict(x_test)
print("Naive Bayes Accuracy with KNN Imputation and with Pregnancy, Insulin, and Skin Thickness Removed", accuracy_score(y_test, y_pred))
print("Confusion Matrix: \n", confusion_matrix(y_test, y_pred))
print("Classification Report Naive Bayes with KNN Imputation and with Pregnancy, Insulin, and Skin Thickness Removed\n", classification_report(y_test, y_pred))
probsnb = nb.predict_proba(x_test)
fprnb, tprnb, thresholds = roc_curve(y_test, probsnb[:, 1], pos_label=1)
roc_aucnb = auc(fprnb, tprnb)

knn = KNeighborsClassifier(n_neighbors = 5, metric = 'minkowski', p = 2)
knn.fit(x_train, y_train)
y_pred = knn.predict(x_test)
print("KNN Accuracy with KNN Imputation and with Pregnancy, Insulin, and Skin Thickness Removed", accuracy_score(y_test, y_pred))
print("Confusion Matrix: \n", confusion_matrix(y_test, y_pred))
print("Classification Report KNN with KNN Imputation and with Pregnancy, Insulin, and Skin Thickness Removed\n", classification_report(y_test, y_pred))
probsknn = knn.predict_proba(x_test)
fprknn, tprknn, thresholds = roc_curve(y_test, probsknn[:, 1], pos_label=1)
roc_aucknn = auc(fprknn, tprknn)

svm = SVC(kernel='linear', probability=True)
svm.fit(x_train, y_train)
y_pred = svm.predict(x_test)
print("SVM Accuracy with Median Imputation and with Pregnancy, Insulin, and Skin Thickness Removed", accuracy_score(y_test, y_pred))
print("Confusion Matrix: \n", confusion_matrix(y_test, y_pred))
print("Classification Report SVM with Median Imputation and with Pregnancy, Insulin, and Skin Thickness Removed\n", classification_report(y_test, y_pred))
probas_ = svm.fit(x_train, y_train).predict_proba(x_test)
# Compute ROC curve and area the curve
fprsvm, tprsvm, thresholds = roc_curve(y_test, probas_[:, 1])
roc_aucsvm = auc(fprsvm, tprsvm)
```


Performance evaluation

Original Code

```
#adaboost = AdaBoostClassifier(random_state=96)
adaboost = AdaBoostClassifier(random_state=96, algorithm='SAMME')
adaboost.fit(x_train, y_train)
y_pred = adaboost.predict(x_test)
print("AdaBoost Accuracy with Median Imputation and with Pregnancy, Insulin, and Skin Thickness Removed", accuracy_score(y_test, y_pred))
print("Confusion Matrix: \n", confusion_matrix(y_test, y_pred))
print("Classification Report AdaBoost with Median Imputation and with Pregnancy, Insulin, and Skin Thickness Removed\n", classification_report(y_test, y_pred))
probsadaboost = adaboost.predict_proba(x_test)
fpradaboost, tpradaboost, thresholds = roc_curve(y_test, probsadaboost[:, 1], pos_label=1)
roc_aucadaboost = auc(fpradaboost, tpradaboost)

#xgboost = XGBClassifier(use_label_encoder=False, eval_metric='mlogloss')
xgboost = XGBClassifier(eval_metric='mlogloss')
xgboost.fit(x_train, y_train)
y_pred = xgboost.predict(x_test)
print("XGBoost Accuracy with Median Imputation and with Pregnancy, Insulin, and Skin Thickness Removed", accuracy_score(y_test, y_pred))
print("Confusion Matrix: \n", confusion_matrix(y_test, y_pred))
print("Classification Report XGBoost with Median Imputation and with Pregnancy, Insulin, and Skin Thickness Removed\n", classification_report(y_test, y_pred))
probsxgboost = xgboost.predict_proba(x_test)
fprxgboost, tprxgboost, thresholds = roc_curve(y_test, probsxgboost[:, 1], pos_label=1)
roc_aucxgboost = auc(fprxgboost, tprxgboost)

fig, ax = plt.subplots(figsize=(7.5, 7.5))

plt.plot(fprdt, tprdt, label='ROC Curve DT Median Imputation Only (AUC = %0.2f)' % (roc_aucdt))
plt.plot(fprlr, tprlr, label='ROC Curve LR Median Imputation Only (AUC = %0.2f)' % (roc_auclr))
plt.plot(fprrf, tprrf, label='ROC Curve RF Median Imputation Only (AUC = %0.2f)' % (roc_aucrf))
plt.plot(fprnb, tprnb, label='ROC Curve NB Median Imputation Only (AUC = %0.2f)' % (roc_aucnb))
plt.plot(fprknn, tprknn, label='ROC Curve KNN Median Imputation Only (AUC = %0.2f)' % (roc_aucknn))
plt.plot(fprsvm, tprsvm, label='ROC Curve SVM Median Imputation Only (AUC = %0.2f)' % (roc_aucsvm))
plt.plot(fpradaboost, tpradaboost, label='ROC Curve Adaboost Median Imputation Only (AUC = %0.2f)' % (roc_aucadaboost))
plt.plot(fprxgboost, tprxgboost, label='ROC Curve XGBoost Median Imputation Only (AUC = %0.2f)' % (roc_aucxgboost))
plt.plot([0, 1], [0, 1], linestyle='--', color='red', label='Random Classifier')
plt.plot([0, 0, 1], [0, 1, 1], linestyle=':', color='green', label='Perfect Classifier')
plt.xlim([-0.05, 1.05])
plt.ylim([-0.05, 1.05])
plt.xlabel('False positive rate')
plt.ylabel('True positive rate')
plt.legend(loc="lower right")
plt.show()
```

Performance evaluation

Refactored Code

Changes made

- Consolidate Duplicate Code
 - Dictionary
 - map training and evaluating each model to its corresponding action, reducing repetitive lines of codes that do the same function.
 - to map print functions

```
yMedian = XtransMedian['target']
xMedianMedian = XtransMedian.drop(['target'], axis=1)
x_train, x_test, y_train, y_test = train_test_split(xMedianMedian, yMedian, test_size=0.30, random_state=14)

models = {
    'Decision Tree': DecisionTreeClassifier(random_state=14),
    'Logistic Regression': LogisticRegression(C=50),
    'Random Forest': RandomForestClassifier(),
    'Naive Bayes': GaussianNB(),
    'KNN': KNeighborsClassifier(n_neighbors=5, metric='minkowski', p=2),
    'SVM': SVC(kernel='linear', probability=True),
    'AdaBoost': AdaBoostClassifier(random_state=96, algorithm='SAMME'),
    'XGBoost': XGBClassifier(eval_metric='mlogloss')
}

models_results = {}
model_phrases = {
    'Decision Tree': {
        'accuracy': "Decision Tree Accuracy with Median Imputation on 10 predictors",
        'classification_report': "Classification Report Decision Tree with Median Imputation on 10 predictors"
    },
    'Logistic Regression': {
        'accuracy': "Logistic Regression Accuracy with Median Imputation and with Pregnancy, Insulin, and Skin Thickness Removed",
        'classification_report': "Classification Report Logistic Regression with Median Imputation and with Pregnancy, Insulin, and Skin Thickness Removed"
    },
    'Random Forest': {
        'accuracy': "Random Forest Accuracy with Median Imputation and with Pregnancy, Insulin, and Skin Thickness Removed",
        'classification_report': "Classification Report Random Forest with Median Imputation and with Pregnancy, Insulin, and Skin Thickness Removed"
    },
    'Naive Bayes': {
        'accuracy': "Naive Bayes Accuracy with KNN Imputation and with Pregnancy, Insulin, and Skin Thickness Removed",
        'classification_report': "Classification Report Naive Bayes with KNN Imputation and with Pregnancy, Insulin, and Skin Thickness Removed"
    },
    'KNN': {
        'accuracy': "KNN Accuracy with KNN Imputation and with Pregnancy, Insulin, and Skin Thickness Removed",
        'classification_report': "Classification Report KNN with KNN Imputation and with Pregnancy, Insulin, and Skin Thickness Removed"
    },
    'SVM': {
        'accuracy': "SVM Accuracy with Median Imputation and with Pregnancy, Insulin, and Skin Thickness Removed",
        'classification_report': "Classification Report SVM with Median Imputation and with Pregnancy, Insulin, and Skin Thickness Removed"
    },
    'AdaBoost': {
        'accuracy': "AdaBoost Accuracy with Median Imputation and with Pregnancy, Insulin, and Skin Thickness Removed",
        'classification_report': "Classification Report AdaBoost with Median Imputation and with Pregnancy, Insulin, and Skin Thickness Removed"
    },
    'XGBoost': {
        'accuracy': "XGBoost Accuracy with Median Imputation and with Pregnancy, Insulin, and Skin Thickness Removed",
        'classification_report': "Classification Report XGBoost with Median Imputation and with Pregnancy, Insulin, and Skin Thickness Removed"
    }
}
```

Performance evaluation

Refactored Code

Changes made

- Consolidate Duplicate Code
 - Used a function to implement Plotting of ROC Curve into one

```
def train_and_evaluate(model, x_train, x_test, y_train, y_test):
    model.fit(x_train, y_train)
    y_pred = model.predict(x_test)
    accuracy = accuracy_score(y_test, y_pred)
    conf_matrix = confusion_matrix(y_test, y_pred)
    class_report = classification_report(y_test, y_pred)
    probas = model.predict_proba(x_test) if hasattr(model, 'predict_proba') else None

    if probas is not None:
        fpr, tpr, _ = roc_curve(y_test, probas[:, 1])
        auc_score = auc(fpr, tpr)
    else:
        fpr, tpr, auc_score = None, None, None

    return accuracy, conf_matrix, class_report, fpr, tpr, auc_score

def print_model_performance(model_name, accuracy, conf_matrix, class_report):
    accuracy_phrase = model_phrases[model_name]['accuracy']
    classification_report_phrase = model_phrases[model_name]['classification_report']

    print(f"{accuracy_phrase}: {accuracy:.4f}")
    print(f"Confusion Matrix: \n{conf_matrix}")
    print(f"{classification_report_phrase}: \n{class_report}")

def plot_roc_curves(models_results):
    fig, ax = plt.subplots(figsize=(7.5, 7.5))

    for model_name, (fpr, tpr, auc_score) in models_results.items():
        if fpr is not None and tpr is not None:
            label = f'ROC Curve {model_name} Median Imputation Only (AUC = %0.2f)' % auc_score
            ax.plot(fpr, tpr, label=label)

    ax.plot([0, 1], [0, 1], linestyle='--', color='red', label='Random Classifier')
    ax.plot([0, 0, 1], [0, 1, 1], linestyle=':', color='green', label='Perfect Classifier')
    ax.set_xlim([-0.05, 1.05])
    ax.set_ylim([-0.05, 1.05])
    ax.set_xlabel('False Positive Rate')
    ax.set_ylabel('True Positive Rate')
    ax.legend(loc="lower right")
    plt.show()

for model_name, model in models.items():
    accuracy, conf_matrix, class_report, fpr, tpr, auc_score = train_and_evaluate(model, x_train, x_test, y_train, y_test)
    print_model_performance(model_name, accuracy, conf_matrix, class_report)
    models_results[model_name] = (fpr, tpr, auc_score)

plot_roc_curves(models_results)
```

Performance evaluation

Refactored Code

Decision Tree Accuracy with Median Imputation on 10 predictors: 0.7355

Confusion Matrix:

```
[[ 97  38]
 [ 35 106]]
```

Classification Report Decision Tree with Median Imputation on 10 predictors:

	precision	recall	f1-score	support
0	0.73	0.72	0.73	135
1	0.74	0.75	0.74	141
accuracy			0.74	276
macro avg	0.74	0.74	0.74	276
weighted avg	0.74	0.74	0.74	276

Logistic Regression Accuracy with Median Imputation and with Pregnancy, Insulin, and Skin Thickness Removed: 0.8297

Confusion Matrix:

```
[[103  32]
 [ 15 126]]
```

Classification Report Logistic Regression with Median Imputation and with Pregnancy, Insulin, and Skin Thickness Removed:

	precision	recall	f1-score	support
0	0.87	0.76	0.81	135
1	0.80	0.89	0.84	141
accuracy			0.83	276
macro avg	0.84	0.83	0.83	276
weighted avg	0.83	0.83	0.83	276

Random Forest Accuracy with Median Imputation and with Pregnancy, Insulin, and Skin Thickness Removed: 0.8406

Confusion Matrix:

```
[[104  31]
 [ 13 128]]
```

Classification Report Random Forest with Median Imputation and with Pregnancy, Insulin, and Skin Thickness Removed:

	precision	recall	f1-score	support
0	0.89	0.77	0.83	135
1	0.81	0.91	0.85	141
accuracy			0.84	276
macro avg	0.85	0.84	0.84	276
weighted avg	0.85	0.84	0.84	276

Performance evaluation

Refactored Code

```
Naive Bayes Accuracy with KNN Imputation and with Pregnancy, Insulin, and Skin Thickness Removed: 0.8152
Confusion Matrix:
[[104  31]
 [ 20 121]]
Classification Report Naive Bayes with KNN Imputation and with Pregnancy, Insulin, and Skin Thickness Removed:
              precision    recall  f1-score   support

      0       0.84       0.77       0.80       135
      1       0.80       0.86       0.83       141

   accuracy          0.82       0.81       0.81       276
  macro avg          0.82       0.81       0.81       276
weighted avg          0.82       0.82       0.81       276

KNN Accuracy with KNN Imputation and with Pregnancy, Insulin, and Skin Thickness Removed: 0.8152
Confusion Matrix:
[[101  34]
 [ 17 124]]
Classification Report KNN with KNN Imputation and with Pregnancy, Insulin, and Skin Thickness Removed:
              precision    recall  f1-score   support

      0       0.86       0.75       0.80       135
      1       0.78       0.88       0.83       141

   accuracy          0.82       0.81       0.81       276
  macro avg          0.82       0.81       0.81       276
weighted avg          0.82       0.82       0.81       276

SVM Accuracy with Median Imputation and with Pregnancy, Insulin, and Skin Thickness Removed: 0.8370
Confusion Matrix:
[[104  31]
 [ 14 127]]
Classification Report SVM with Median Imputation and with Pregnancy, Insulin, and Skin Thickness Removed:
              precision    recall  f1-score   support

      0       0.88       0.77       0.82       135
      1       0.80       0.90       0.85       141

   accuracy          0.84       0.84       0.84       276
  macro avg          0.84       0.84       0.84       276
weighted avg          0.84       0.84       0.84       276
```

Performance evaluation

Refactored Code

```
AdaBoost Accuracy with Median Imputation and with Pregnancy, Insulin, and Skin Thickness Removed: 0.8333
Confusion Matrix:
[[106  29]
 [ 17 124]]
Classification Report AdaBoost with Median Imputation and with Pregnancy, Insulin, and Skin Thickness Removed:
              precision    recall  f1-score   support

      0       0.86       0.79       0.82       135
      1       0.81       0.88       0.84       141

   accuracy          0.83       276
  macro avg       0.84       0.83       0.83       276
weighted avg       0.84       0.83       0.83       276

XGBoost Accuracy with Median Imputation and with Pregnancy, Insulin, and Skin Thickness Removed: 0.8188
Confusion Matrix:
[[105  30]
 [ 20 121]]
Classification Report XGBoost with Median Imputation and with Pregnancy, Insulin, and Skin Thickness Removed:
              precision    recall  f1-score   support

      0       0.84       0.78       0.81       135
      1       0.80       0.86       0.83       141

   accuracy          0.82       276
  macro avg       0.82       0.82       0.82       276
weighted avg       0.82       0.82       0.82       276
```

Performance evaluation

Refactored Code

